



Predictive Modeling of Health Outcomes Using Blood Glucose Time- Series Data



BDA Complex Computing Problem Detailed Report
SANAULLAH TAREEN | M. Mansoor Khan | Aitsam Zulfiqar

Detailed CCP Report

Predictive Modeling of Health Outcomes Using Blood Glucose Time-Series Data

1. Introduction

This project presents a predictive modeling framework for blood glucose time-series data in a healthcare setting. The goal is to forecast future glucose values and identify clinically meaningful high-risk conditions such as hyperglycemia and hypoglycemia using machine learning. The work is positioned within the broader domain of health analytics and predictive modeling, which aligns with the CCP focus on computational solutions for medical data.

The study emphasizes time-dependent prediction, model training, hyperparameter tuning, and clinical interpretation. In addition, the notebook includes a safety-oriented alert logic that evaluates whether the model output supports appropriate clinical action. This makes the work relevant to decision-support scenarios where prediction quality directly affects patient care.

2. CCP Context and Fit

The CCP brief requires a health-related predictive modeling problem based on longitudinal data, with clear outcome definition, feature engineering, model evaluation, and interpretability. It also highlights the importance of clinically meaningful outcomes, use of appropriate metrics, and strong justification of algorithmic choices. This project addresses those expectations through a blood glucose forecasting task with clinical alert validation.

Blood glucose is a core physiological variable in diabetes management and inpatient monitoring. Predicting future glucose levels is not only a forecasting problem, but also a practical healthcare analytics problem because it supports early detection of dangerous trends. The project therefore fits the CCP theme of using computational methods to support health outcome prediction.

3. Problem Statement

The objective of this project is to predict future blood glucose values from time-series data and determine whether the predicted value indicates a clinically important state. If the forecast suggests hyperglycemia, the system should support insulin reduction guidance; if it suggests hypoglycemia, it should support insulin escalation guidance. This transforms a raw regression task into a clinically grounded prediction workflow.

The problem is computationally meaningful because it combines sequence learning, regression modeling, model optimization, and clinical decision logic. It also involves handling time-dependent data where recent patterns influence future risk. Such structure is consistent with the CCP requirement for a non-trivial predictive modeling problem in healthcare analytics.

4. Data and Features

The notebook is built around blood glucose time-series information, which represents a longitudinal medical signal. Time-series data are especially important in healthcare because physiological states change over time rather than remaining static. The modeling process therefore captures dynamic patterns that are essential for predicting future values.

The project can be interpreted as a patient monitoring task where the target variable is the next blood glucose value or future glucose trend. In a CCP-aligned framing, glucose history serves as the core physiological sequence, while additional patient-level variables such as age, sex, diabetes status, medications, or comorbidities can be incorporated to strengthen the clinical relevance. This supports the CCP requirement for mixed feature types and longitudinal analysis.

5. Methodology

The notebook uses CatBoost as the main predictive model and includes hyperparameter tuning to improve performance. CatBoost is a strong choice for healthcare prediction tasks because it handles non-linear relationships well and is effective in settings with structured features. The use of optimization demonstrates that model selection was not arbitrary, but based on systematic search for a stronger configuration.

The workflow also includes train-test evaluation and reported error values, which are essential for predictive modeling. The notebook shows the best CatBoost configuration and test performance, indicating that the project has moved beyond a basic prototype into a more disciplined modeling pipeline. This is important for the CCP because the brief asks for model justification, performance evaluation, and implementation quality.

6. Model Training and Results

The notebook reports a train RMSE of 1.6030 and a test RMSE of 1.5608 for CatBoost. These values indicate stable generalization, since the test error is not worse than the training error. The optimized CatBoost search also produced a best RMSE of about 1.3223 during tuning, showing that the

project explored better parameter settings rather than relying on a default model.

The training logs show gradual improvement during optimization and fitting, which suggests that the model is learning useful structure from the data. This supports the claim that the project is not a trivial baseline, but a genuine predictive workflow. For CCP purposes, this kind of result is useful because it demonstrates implementation, experimentation, and performance analysis.

7. Clinical Interpretation

The strongest clinical value of this project comes from its alert logic. The notebook checks predicted blood glucose against thresholds for hyperglycemia and hypoglycemia and then validates whether the suggested action is clinically sensible. This turns the prediction output into a meaningful medical signal rather than just a numeric forecast.

A predicted high glucose value can be treated as a warning for excess glycemic exposure, while a predicted low glucose value can indicate risk of acute hypoglycemia. Such interpretation is useful because the CCP asks for clinically meaningful language and ethical relevance, not just raw model scores. The project therefore supports the idea of decision assistance in monitoring and intervention.

8. Evaluation Against CCP

This project aligns with several CCP requirements. It uses a healthcare-related dataset, applies a machine learning model, performs prediction, includes time-series structure, and introduces a clinically meaningful

validation check. It also demonstrates computational effort through CatBoost training and hyperparameter optimization.

The project can be strengthened further by presenting the glucose forecasting task explicitly as a clinical risk prediction problem. For example, the output can be reframed as future hypoglycemia risk, hyperglycemia risk, or an adverse glycemic event within a prediction window. This would fit even more closely with the CCP's emphasis on predictive modeling of health outcomes.

9. Complexity and Computation

The modeling pipeline is computationally non-trivial because time-series prediction requires capturing temporal dependence in the data. CatBoost adds further complexity by using gradient boosting over multiple trees and by supporting parameter tuning across several model configurations. This makes the work suitable for a big-data analytics assignment rather than a simple classification exercise.

The project also includes practical decision logic after prediction, which increases the number of subproblems being solved. In CCP terms, this reflects interdependence between forecasting, thresholding, and clinical action validation. Such structure supports the assignment's requirement for analytical reasoning and multi-step problem solving.

10. Interpretation of Quality

The project is technically strong because it combines predictive modeling with clinically relevant alert logic. It is also easy to defend because the healthcare use case is clear: monitoring glucose trends and identifying

dangerous excursions. The notebook shows effort in optimization, numerical evaluation, and health-context interpretation.

From a CCP perspective, the report should emphasize the health outcome focus, the predictive framework, and the clinical usefulness of the results. That framing presents the project as a practical and defensible healthcare analytics solution. The main message is that the work is relevant, algorithmically substantial, and aligned with the intent of the CCP.

11. Conclusion

This project provides a solid foundation for a healthcare predictive modeling report under the CCP. It uses blood glucose time-series data, applies a strong regression model, performs optimization, and converts outputs into clinically meaningful alerts. These elements make the project suitable for a health analytics submission.

For the final report, the strongest presentation is to treat the work as a clinical forecasting and risk-support system. With that framing, the project clearly demonstrates machine learning, time-series reasoning, evaluation, and medical interpretation in a way that is consistent with the CCP objectives.